

# Step-by-Step: FastAPI on Kubernetes

## Step 1: Create a minimal FastAPI app

`app/main.py`

```
from fastapi import FastAPI, Request

app = FastAPI()

@app.get("/health")
def health():
    return {"status": "ok"}

@app.post("/minio-event")
async def minio_event(request: Request):
    event = await request.json()
    print(event) # replace with real logic
    return {"status": "received"}
```

`requirements.txt`

```
fastapi
uvicorn[standard]
```

---

## Step 2: Create Dockerfile

`Dockerfile`

```
FROM python:3.11-slim

WORKDIR /app

COPY requirements.txt .
```

```
RUN pip install --no-cache-dir -r requirements.txt  
  
COPY app/ app/  
  
EXPOSE 8000  
  
CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Build & push:

```
docker build -t myrepo/fastapi-minio-webhook:1.0 .
docker push myrepo/fastapi-minio-webhook:1.0
```

---

## Step 3: Kubernetes Deployment

`deployment.yaml`

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: fastapi-webhook
spec:
  replicas: 2
  selector:
    matchLabels:
      app: fastapi-webhook
  template:
    metadata:
      labels:
        app: fastapi-webhook
    spec:
      containers:
        - name: fastapi
          image: myrepo/fastapi-minio-webhook:1.0
          ports:
            - containerPort: 8000
          resources:
```

```
requests:
  cpu: "100m"
  memory: "128Mi"
limits:
  cpu: "500m"
  memory: "512Mi"
readinessProbe:
  httpGet:
    path: /health
    port: 8000
  initialDelaySeconds: 5
  periodSeconds: 10
livenessProbe:
  httpGet:
    path: /health
    port: 8000
  initialDelaySeconds: 15
  periodSeconds: 20
```

Apply:

```
kubectl apply -f deployment.yaml
```

---

## Step 4: Kubernetes Service (ClusterIP)

MinIO will call this **inside the cluster**.

**service.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: fastapi-webhook
spec:
  selector:
    app: fastapi-webhook
  ports:
```

```
- port: 8000
  targetPort: 8000
type: ClusterIP
```

Apply:

```
kubectl apply -f service.yaml
```

Service DNS (important ⚡):

```
http://fastapi-webhook.default.svc.cluster.local:8000
```

---

## Step 5: Test inside the cluster

```
kubectl run curltest --rm -it \
--image=curlimages/curl -- sh
curl http://fastapi-webhook:8000/health
```

Expected:

```
{"status": "ok"}
```

---

## Step 6: Configure MinIO webhook endpoint

In MinIO config:

```
endpoint=http://fastapi-webhook:8000/minio-event
```

If MinIO is in a different namespace:

```
http://fastapi-webhook.<namespace>.svc.cluster.local:8000/minio-event
```

---

## Step 7: Logs & Debugging

```
kubectl get pods  
kubectl logs deploy/fastapi-webhook
```

---

## Optional (Highly Recommended in Prod)

### Add Auth Header Validation

```
from fastapi import Header, HTTPException  
  
@app.post("/minio-event")  
async def minio_event(  
    request: Request,  
    x_token: str = Header(None)  
):  
    if x_token != "super-secret":  
        raise HTTPException(status_code=401)
```

Then set header in MinIO webhook config.

---

### Scale Automatically (HPA)

```
apiVersion: autoscaling/v2  
kind: HorizontalPodAutoscaler  
metadata:  
  name: fastapi-webhook-hpa  
spec:  
  scaleTargetRef:  
    apiVersion: apps/v1  
    kind: Deployment  
    name: fastapi-webhook  
  minReplicas: 2  
  maxReplicas: 10
```

```
metrics:
- type: Resource
  resource:
    name: cpu
  target:
    type: Utilization
    averageUtilization: 70
```

---

## Expose Externally (only if needed)

- Ingress (NGINX / Traefik)
  - NodePort / LoadBalancer (less common for webhooks)
-